

# Ravi\_Chandra\_Reddy\_Basireddy\_CSCI544\_HW1

September 8, 2022

## 1 CSCI544 Homework1 - Ravi Chandra Reddy Basireddy

Python Version : 3.9

### 1.1 Imports

- Pandas: To work with dataframes.
- NLTK: A Natural Language Toolkit used for processing textual data.
- RE: Regular Expressions used for handling word findings & substitutions.
- BS4: BeautifulSoup Library is a parser that can handle HTML Tags and Links.
- Contractions: A library to contract and de-contract Contractions.
- String: A library to handle strings.
- Warning: A library to handle console warnings.

```
[1]: import pandas as pd
import nltk
nltk.download('wordnet',quiet=True)
import re
from bs4 import BeautifulSoup
import contractions
import string
import warnings
warnings.filterwarnings(action='ignore')

#!pip install bs4
#!pip install contractions
#!pip install nltk
#!pip install string
#!pip install pandas
#!pip install warnings
#!pip install sklearn
```

### 1.2 Read Data

- Read Data from a TSV file where the data is separated using tabs.
- we are only interested in Star Rating and Review Body.
- Star Rating: Rating given by the customers in the range of 1 to 5.
- Review Body: Review given by the customers in the textual format.

```
[2]: amazon_reviews=pd.read_csv('data.tsv',
    ↳sep='\t',usecols=['star_rating','review_body'],low_memory=False)
```

### 1.3 Keep Reviews and Ratings

- Already completed at the reading data step.
- Dropping NaN Values which has no meaning to the rating.
- Dropping Duplicates which are repeated.
- Printing out the first five values to know what kind of data is in dataframe.

```
[3]: amazon_reviews=amazon_reviews.dropna()
amazon_reviews=amazon_reviews.drop_duplicates()
amazon_reviews.head(5)
```

```
[3]:  star_rating      review_body
0          5  so beautiful even tho clearly not high end ...
1          5  Great product.. I got this set for my mother, ...
2          5  Exactly as pictured and my daughter's friend l...
3          5  Love it. Fits great. Super comfortable and nea...
4          5  Got this as a Mother's Day gift for my Mom and...
```

### 1.4 We select 20000 reviews randomly from each rating class.

- filtering out the data with respective labels.
- sampling 20k reviews from each class.
- Combining all the data from differnt classes to create a vector of Dimension (100000,2).

```
[4]: star_one=amazon_reviews[amazon_reviews.star_rating=='1']
star_one=star_one.sample(n=20000)
star_two=amazon_reviews[amazon_reviews.star_rating=='2']
star_two=star_two.sample(n=20000)
star_three=amazon_reviews[amazon_reviews.star_rating=='3']
star_three=star_three.sample(n=20000)
star_four=amazon_reviews[amazon_reviews.star_rating=='4']
star_four=star_four.sample(n=20000)
star_five=amazon_reviews[amazon_reviews.star_rating=='5']
star_five=star_five.sample(n=20000)
sampled_reviews=pd.
    ↳concat([star_one,star_two,star_three,star_four,star_five],ignore_index=True)
sampled_reviews.shape
```

```
[4]: (100000, 2)
```

## 2 Data Cleaning

- Cleaning the data inorder to make the models better, as better data will always result better Prediction.

### 3 Pre-processing

1. Removing URL.
2. Removing HTML Tags.
3. Removing All the characters except for A-Z&a-z.
4. Removing any html text left with BeautifulSoup Library.
5. Removing Contractions.
6. Removing Punctuation.
7. Removing extra Spaces.
8. Converting the text to lowercase.

```
[5]: def remove_punctuation(review):  
      return ''.join([words for words in review if words not in string.  
      ↪punctuation ])
```

```
[6]: def clean_review(review):  
      review = re.sub(r"http\S+", "", review)  
      review = re.sub('<.*?>+', '', review)  
      review = re.sub('[^A-Za-z]+', ' ', review)  
      review = BeautifulSoup(review, "html.parser").get_text()  
      review = contractions.fix(review)  
      review = remove_punctuation(review)  
      review = re.sub("\S*\d\S*", "", review).strip()  
      review = review.lower()  
      return review
```

#### 3.0.1 Calculating the Average Length of Reviews by Character

```
[7]: def average_count(sampled_reviews):  
      number_of_sentences=len(sampled_reviews)  
      return sum(map(len,sampled_reviews))/number_of_sentences
```

```
[8]: beforeCleaning=average_count(sampled_reviews['review_body'])  
      sampled_reviews['review_body']=sampled_reviews['review_body'].apply(lambda_  
      ↪review:clean_review(review))  
      afterCleaning=average_count(sampled_reviews['review_body'])  
      print("Average character length of the reviews Before and After_  
      ↪Cleaning",beforeCleaning,',',afterCleaning)
```

Average character length of the reviews Before and After Cleaning 198.21343 ,  
190.29407

#### 3.1 remove the stop words

Removing the stop words such as “The” which have no meaning to them.

```
[9]: from nltk.corpus import stopwords  
      stop_words = set(stopwords.words('english'))
```

```

beforePreprocessing=average_count(sampled_reviews['review_body'])
sampled_reviews['review_body']= sampled_reviews['review_body'].apply(lambda
    ↪review: " ".join([word for word in review.split() if word not in
    ↪stop_words]))

```

### 3.2 perform lemmatization

Converting all the words in the reviews to single form so that they can be matched by similarity.

```

[10]: from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
sampled_reviews['review_body']= sampled_reviews['review_body'].apply(lambda
    ↪review: ' '.join(lemmatizer.lemmatize(e) for e in nltk.
    ↪word_tokenize(review)))
afterPreprocessing=average_count(sampled_reviews['review_body'])
print("Average character length of the reviews Before and After
    ↪Preprocessing",beforePreprocessing,',',afterPreprocessing)

```

Average character length of the reviews Before and After Preprocessing 190.29407  
, 112.67224

## 4 TF-IDF Feature Extraction

- Converting Reviews to Count Vectors using a concept known as TF-IDF.
- It is the relation between Term Frequency and Inverse Document Frequency.
- Term Frequency is the frequency of the word in a corpus.
- Inverse Document Frequency is the Frequency of a Word in that particular Document.
- TF - IDF tells about the Frequency of a Word in that Particular Document With Respect to the Entire Corpus.
- N-Grams are combination of words in that particular document. Bi-gram Example (really-appreciate).

```

[11]: from sklearn.feature_extraction.text import TfidfVectorizer
tf_idf_vect=TfidfVectorizer(ngram_range=(1,3))
final_tf_idf=tf_idf_vect.fit_transform(sampled_reviews['review_body'].values)

```

### 4.1 Train Test Split

- We split the data in the split of 80:20 which is 80% of the for Training and 20% of the Data for Testing.
- We use train test split in order to train the model and test performance of the model.

```

[12]: from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(final_tf_idf,
    ↪sampled_reviews['star_rating'], test_size = 0.2)

```

## 4.2 Classification Metrics

A Function that gives you information on Accuracy, Precision, Recall and F1-score.

```
[13]: from sklearn.metrics import classification_report, confusion_matrix, \
      ↪ accuracy_score

def metrics(prediction, actual):
    print('\nAccuracy:', accuracy_score(actual, prediction))
    print('\nclassification_report\n')
    print(classification_report(actual, prediction))
```

## 5 Perceptron

- Perceptron is a two class classification Model.
- It uses a concept of Neuron, which has an activation function, which activates only when crossing a certain threshold.
- We used `random_state` to randomize the data.
- We used `n_jobs` to Run Parallel on All Cores.
- The Accuracy of this model on this data is in the range of 45-47.

```
[14]: from sklearn.linear_model import Perceptron
perceptronModel = Perceptron(random_state=0,n_jobs=-1)
perceptronModel.fit(xtrain, ytrain)
predictions=perceptronModel.predict(xtest)
metrics(predictions, ytest)
```

Accuracy: 0.45805

classification\_report

	precision	recall	f1-score	support
1	0.56	0.55	0.56	3965
2	0.36	0.35	0.36	4051
3	0.37	0.30	0.33	4044
4	0.40	0.41	0.41	3932
5	0.55	0.69	0.61	4008
accuracy			0.46	20000
macro avg	0.45	0.46	0.45	20000
weighted avg	0.45	0.46	0.45	20000

## 6 SVM

- SVM uses a concept of boundary, which helps it to detect and avoid outliers.

- SVM have different form of Kernel: Linear, Poly and more, which can be used for different forms of data.
- We used C=0.1 which is a regularization parameter.
- The Accuracy of this model on this data is in the range of 50-52.

```
[15]: from sklearn.svm import LinearSVC
SVM = LinearSVC(C=0.1)
SVM.fit(xtrain, ytrain)
predictions=SVM.predict(xtest)
metrics(predictions, ytest)
```

Accuracy: 0.51835

classification\_report

	precision	recall	f1-score	support
1	0.54	0.71	0.62	3965
2	0.43	0.31	0.36	4051
3	0.46	0.37	0.41	4044
4	0.49	0.41	0.45	3932
5	0.60	0.79	0.68	4008
accuracy			0.52	20000
macro avg	0.50	0.52	0.50	20000
weighted avg	0.50	0.52	0.50	20000

## 7 Logistic Regression

- Logistic Regression tells the likelihood between the classes.
- Logistic regression uses logarithms to compress the data between 0 and 1 which act similar to probability.
- We used solver as Saga which is fastest and best for huge data. There are other solvers like newton-cg, lbfgs & more.
- We used random\_state to randomize the data.
- We used n\_jobs to Run Parallel on All Cores.
- We used max iterations as 200.
- The Accuracy of this model on this data is in the range of 51-53.

```
[16]: from sklearn.linear_model import LogisticRegression
logisticModel=LogisticRegression(solver='saga',random_state=0,n_jobs=-1,max_iter=200)
logisticModel.fit(xtrain, ytrain)
predictions=logisticModel.predict(xtest)
metrics(predictions, ytest)
```

Accuracy: 0.52425

classification\_report

	precision	recall	f1-score	support
1	0.57	0.65	0.61	3965
2	0.42	0.37	0.40	4051
3	0.44	0.42	0.43	4044
4	0.49	0.45	0.47	3932
5	0.65	0.73	0.69	4008
accuracy			0.52	20000
macro avg	0.52	0.52	0.52	20000
weighted avg	0.52	0.52	0.52	20000

## 8 Naive Bayes

- Naive Bayes uses the principle of Bayes Theorem.
- Bayes Theorem makes use of Conditional Probability.
- Naive Bayes is one of the fastest as it computes probabilities which require little to none computing.
- The Accuracy of this model on this data is in the range of 51-53.

```
[17]: from sklearn import naive_bayes
naiveBayesModel = naive_bayes.MultinomialNB()
naiveBayesModel.fit(xtrain, ytrain)
predictions=naiveBayesModel.predict(xtest)
metrics(predictions, ytest)
```

Accuracy: 0.5143

classification\_report

	precision	recall	f1-score	support
1	0.59	0.64	0.61	3965
2	0.42	0.36	0.39	4051
3	0.44	0.36	0.39	4044
4	0.44	0.50	0.47	3932
5	0.65	0.71	0.68	4008
accuracy			0.51	20000
macro avg	0.51	0.52	0.51	20000
weighted avg	0.51	0.51	0.51	20000